

{Kariera Ruby Coding School - June 2017}

# DATA MODELLING

AND RELATIONAL DATABASES

A g e n t l e i n t r o d u c t i o n

Dr. John Pagonis

@zanshinlabs



**Zanshin**  
Labs

**THE DONUTS ARE  
SPONSORED BY...**



**Zanshin**  
Labs





**software production**  
consultants



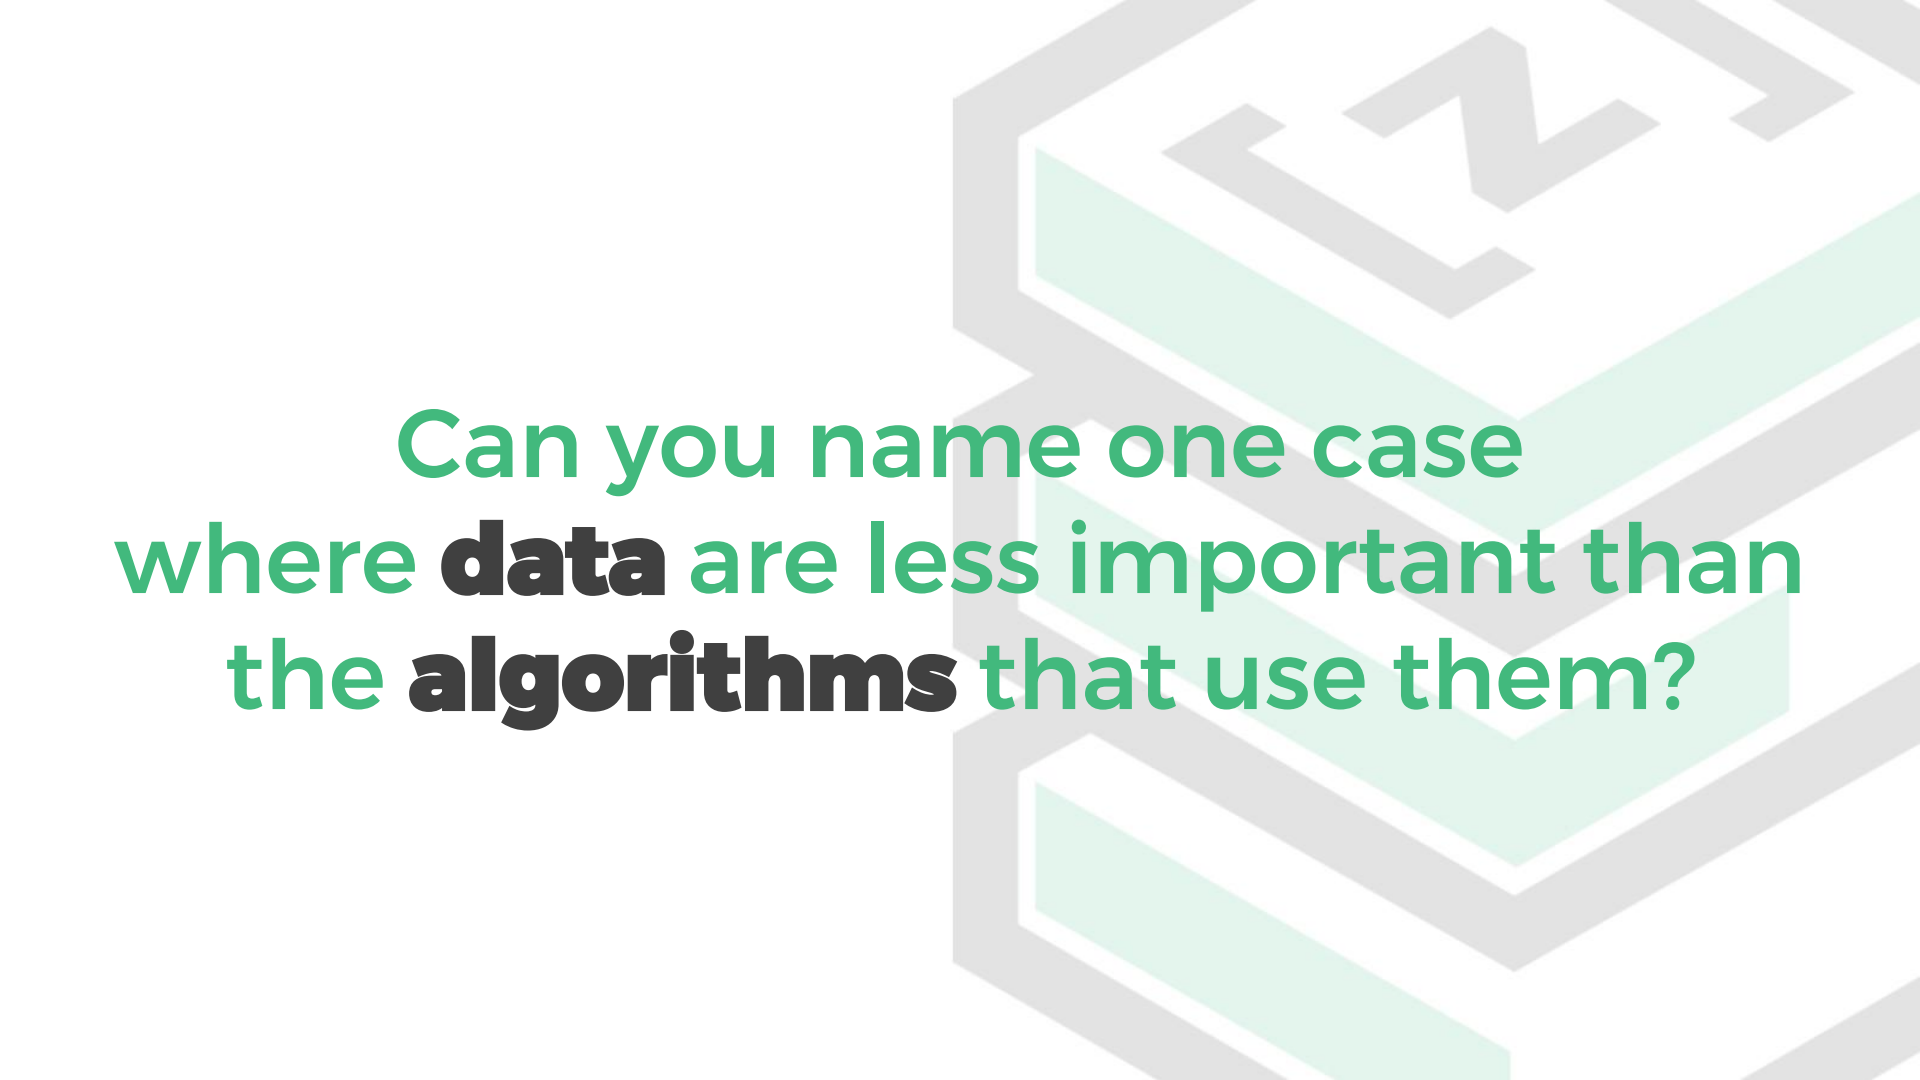
we help teams & organisations

**define** | **design** | **develop**

software their **users need**



# **A STORY ABOUT DATA**



Can you name one case  
where **data** are less important than  
the **algorithms** that use them?

Ancient computing proverb....

# Garbage **In** Garbage **Out**

---

It doesn't matter how smart your algorithm is when you  
feed it with garbage....

Ancient computing proverb....


# Garbage **In** Garbage **Out**

---

This also applies for your shiny new Rails app...



This is your app



---

Make sure that the kernel of your app is not rotten!!

What about Rails?

Rails is a Ruby framework for  
building **data(base) driven** web  
applications

It's business as usual

Organisations need (IT) to  
perform their day-to-day  
**operations** and also make  
**decisions**

---

It has always been like this since the ancient times.

It's business as usual

Organisations need (IT) to  
perform their day-to-day  
**operations** and also make  
**decisions**

---

To achieve this, they need **information**.

The background features a complex geometric pattern. On the right side, there are several stacked, isometric blocks in shades of light green and grey. To the left of these blocks, a grey maze-like structure is visible, consisting of interconnected lines and angles. The overall design is modern and technical.

What is the difference between  
**data and information?**

## Data vs information

Data is the collection of **raw facts**,  
whereas information is those data  
transformed in an **actionable** and  
**meaningful** form

---

i.e. Can you learn something? Can you do something?

The background features a complex geometric pattern. On the right side, there are several stacked, isometric blocks in shades of light green and grey. To the left of these blocks, a grey maze-like structure is visible, with a path leading towards the center. The overall design is clean and modern, with a focus on geometric shapes and a muted color palette.

How can we transform **data** into  
**information?**



To transform **data**  
to **information**

Combine





To transform **data**  
to **information**

Combine  
Compare



# To transform **data** to **information**

Combine  
Compare  
Filter



# To transform **data** to **information**

Combine  
Compare  
Filter  
Summarise



# To transform **data** to **information**

Combine  
Compare  
Filter  
Summarise  
Organise

About transforming data to information

The more **flexible** you are in doing  
such transformations the **better**

---

The **more ways** you see data the **more info** you get!

About transforming data to information

The more **accurate** you are in doing  
such transformations the **better**

---

G.I.G.O

About transforming data to information

The **fastest** you are in doing  
such transformations the **better**

---

e.g., for event based decisions

It's business as usual

Organisations need information  
to run their **operations**  
and make **decisions**


---

Provide and capture the right information to do the  
**job right.**





**How powerful**  
(flexible, fast, accurate)  
do you want to be at storing,  
retrieving and producing  
information?



# **ABOUT DATABASES**



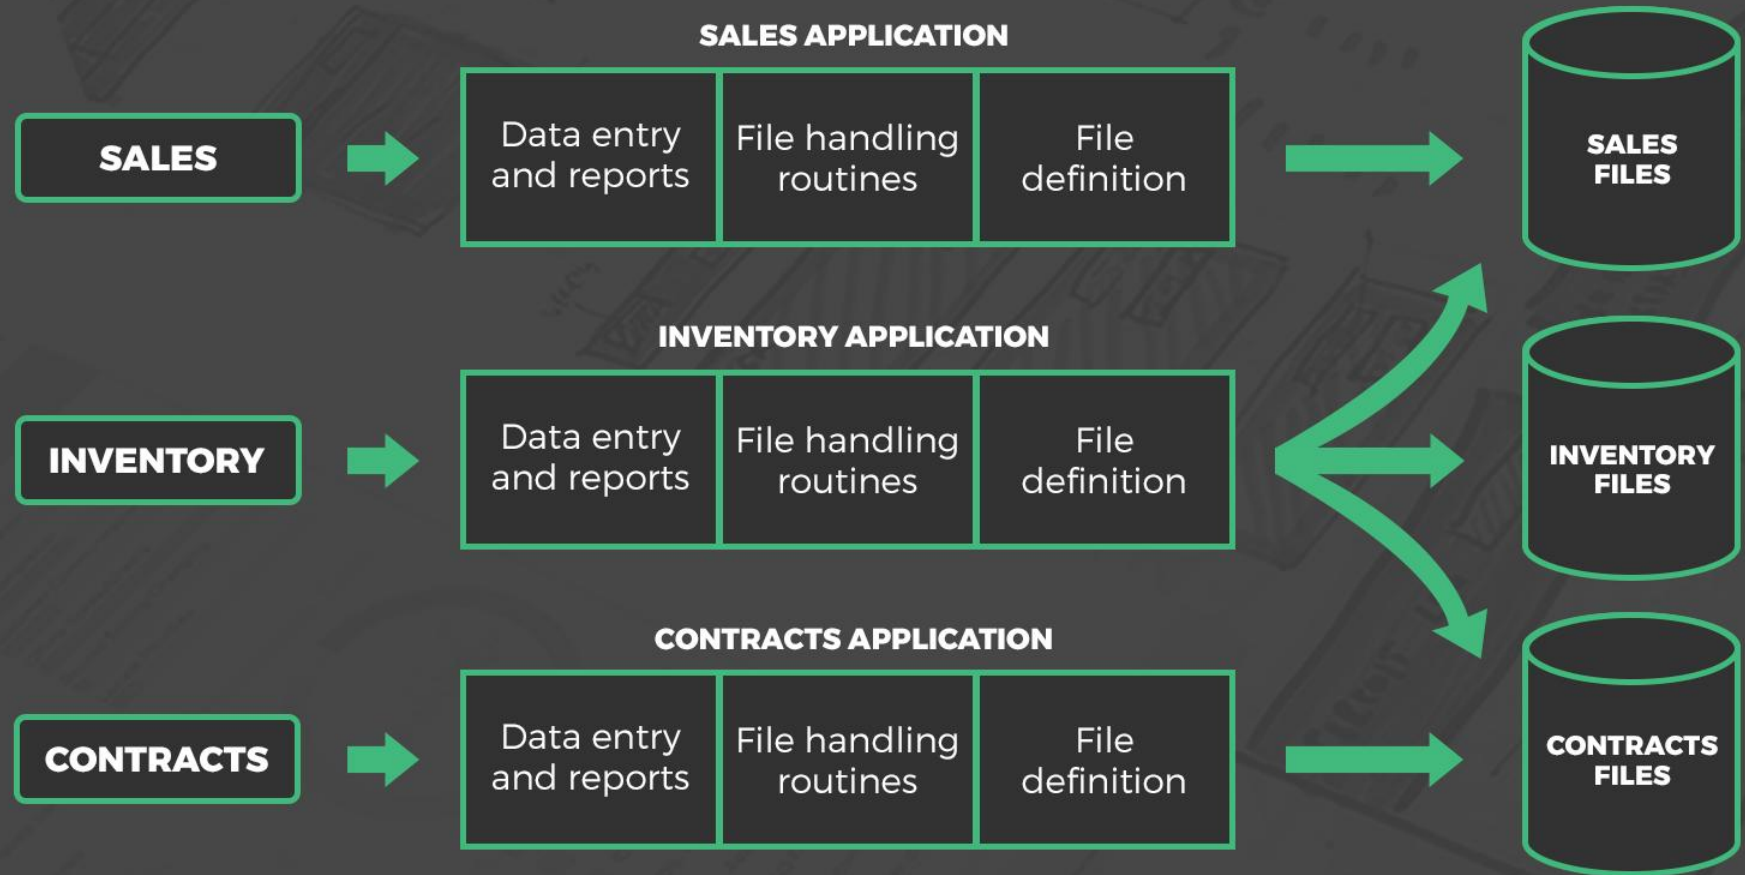
What used to happen  
(at organisations) before  
**databases?**

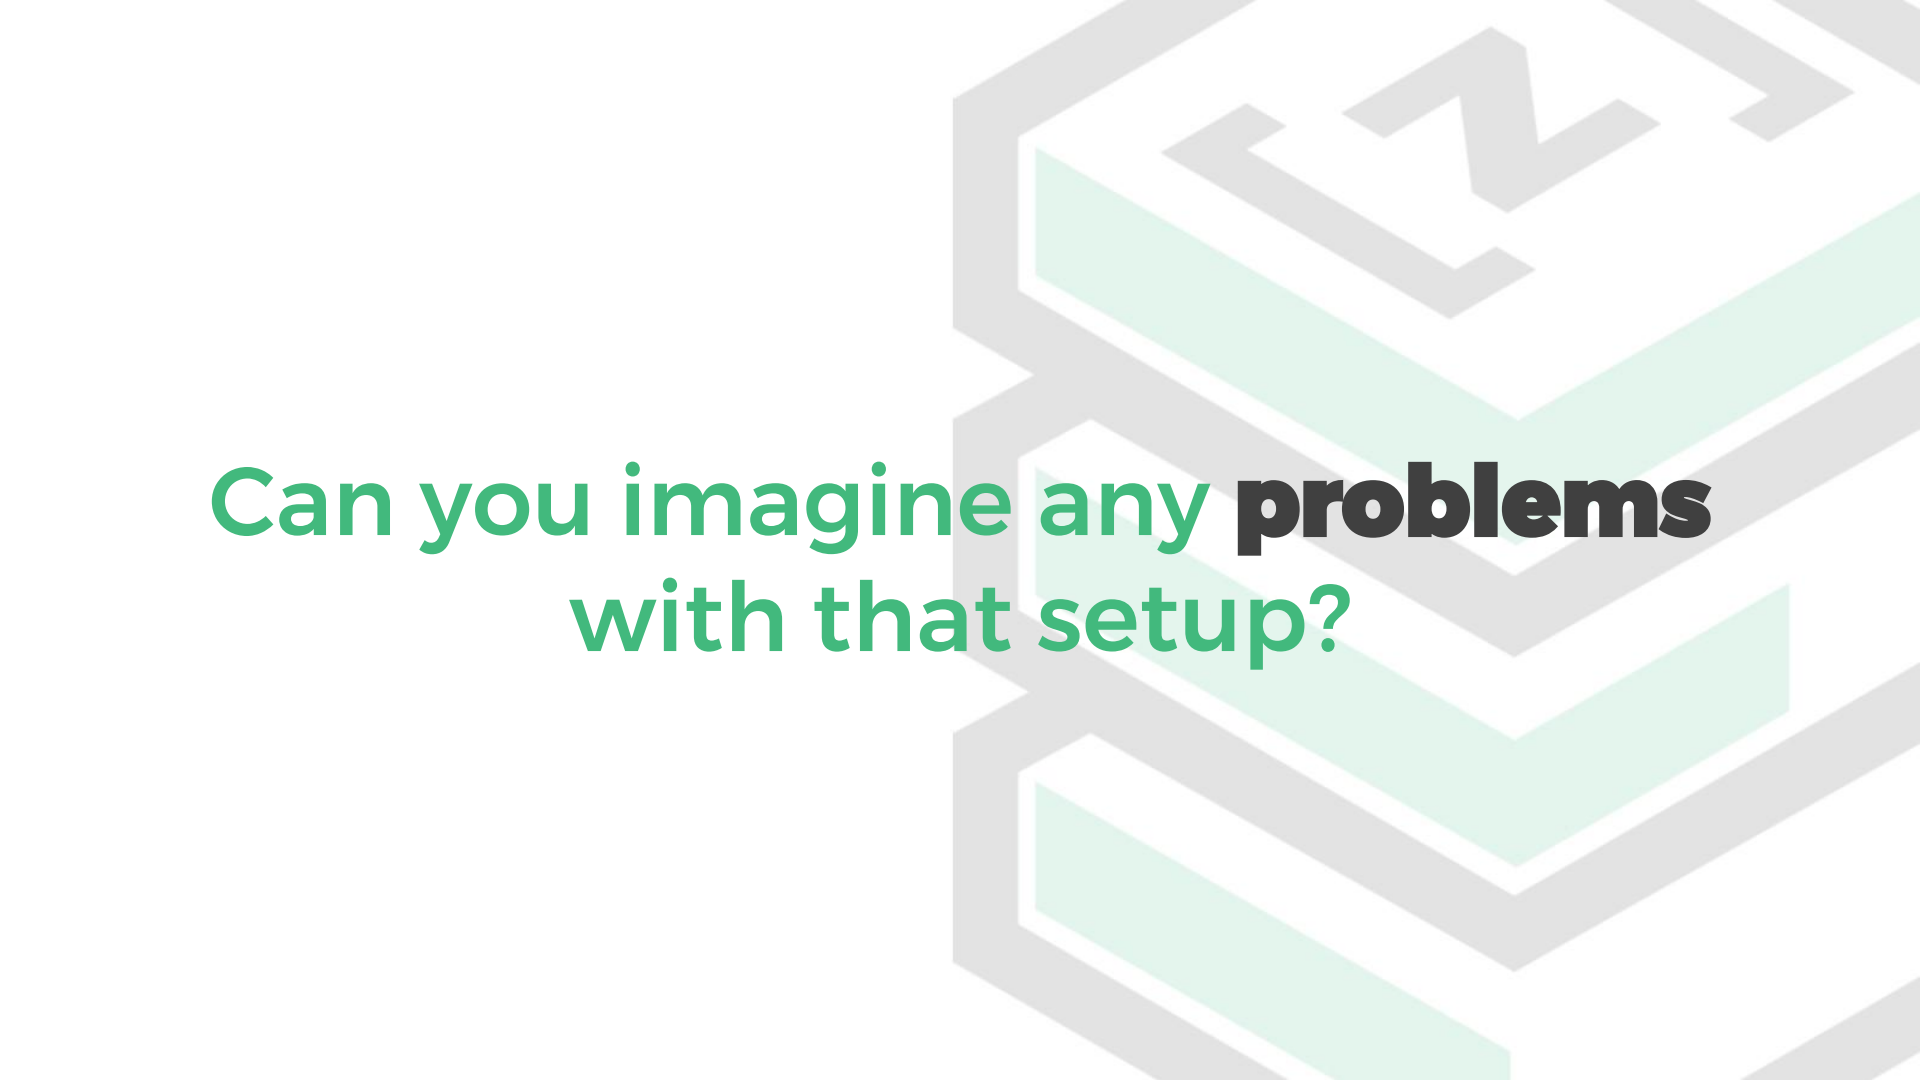
File-based processing...

We had **flat files** being accessed  
by many programs **directly**.

Each one of these had to know  
the exact **same rules** about the  
business and the data.

## File-based processing...



The background features a complex geometric pattern. On the right side, there are several stacked, isometric blocks in shades of light green and grey. To the left of these blocks, a grey maze-like structure is visible, consisting of interconnected lines forming a path. The overall design is clean and modern, with a focus on geometric shapes and patterns.

Can you imagine any **problems**  
with that setup?

The background features a series of three-dimensional, isometric blocks stacked in a descending staircase pattern from the top right towards the bottom left. The blocks are colored in a light mint green and a pale grey. A large, stylized, grey 'M' shape is positioned on the top-most block, appearing as if it's a logo or a decorative element. The overall aesthetic is clean, modern, and geometric.

So what are **databases** then?

Collection of logically **related data** with descriptions about these data (i.e. **metadata**)

---

Metadata in databases provide **program-data independence** (among other things).



# Databases

Collection of logically **related data** with descriptions about these data (i.e. **metadata**)

---

In a database, conceptually (and many times physically) everything is stored **in one place**.

# DataBase Management Systems

A DB is typically directly accessed  
by one program, the **DBMS**

---

The DBMS **enforces constraints** to the database.



So, is there a major difference  
between **MySQL** and **SQLite**?

## DataBase Management Systems

A **client/server** DB is only directly accessed by one program, the **DBMS**.  
It is the **gateway** through which other programs access the data

---

The DBMS **enforces constraints** to the database.

A taxonomy of databases (there are many taxonomies)

Hierarchical  
Relational  
Document-oriented  
Object-oriented  
Object-relational  
Graph

# **RELATIONAL DATABASES**



Relational databases (strictly speaking)

A relational database is one  
which is based on the relational  
model of data as defined by  
**E.F Codd**

---

Codd, E.F. (1970). "A Relational Model of Data for Large Shared Data Banks". Communications of the ACM. 13 (6): 377-387.

## Relational databases

A relational database stores data about multiple different **entities**

---

Entities are stored in **tables**.



## Relational databases

A relational database stores data about the **relationships** between these **entities**

---

The relationships are **links** between tables.

Relational databases (in simple terms)

A relational database is a collection of **entities** and the **relationships** between them

---

Relational databases (in simple terms)

A relational database is a collection of **entities** and the **relationships** between them

---

Each **table** should be about **one entity**.  
Everything should only be there **once**.

Relational databases (in simple terms)

A table is a collection of data about one **entity**. It is a collection of **rows**. Each row has the same **columns**.

---

An **entity** is typically expressed and identified as a **noun**.  
If a table has data about more than one entity, then it has **redundant** data.

Relational databases (in simple terms)

A column is an **attribute** of an entity. It is one elemental piece of data.

---

An **attribute** may also be expressed as a **noun**.

Relational databases (in simple terms)

A row is a collection of columns  
about **one member** of an entity.

---

An **attribute** may also be expressed as a **noun**.

Relational databases (in simple terms)

A primary key is a special  
column (or set of)  
that **uniquely identifies** a row.

---

Relational databases (in simple terms)

A primary key is a special column (or set of) that **uniquely identifies** a row. There can only be one per table. It can't be null.

---



Relational databases (in simple terms)

A primary key is a special column (or set of) that **uniquely identifies** a row. There can only be one per table. It can't be null.

---

If the entity has one then every row **has one**.

Relational databases (in simple terms)

A primary key is a special column (or set of) that **uniquely identifies** a row. There can only be one per table. It can't be null.

---

If the entity has one then every row **has one**.  
No two rows have the same one.

Relational databases (in simple terms)

A primary key is a special column (or set of) that **uniquely identifies** a row. There can only be one per table. It can't be null.

---

If the entity has one then every row **has one**.  
No two rows have the same one.  
There are **natural** and **surrogate** primary keys.

## Relational databases (in simple terms)

A foreign key is an attribute in one table that **refers** to a primary key in another table.

---

The background features a series of overlapping, three-dimensional geometric shapes. These shapes are primarily light green and light grey, creating a sense of depth and perspective. The shapes are arranged in a way that suggests a stack of blocks or a series of steps, with some shapes having a more complex, angular design. The overall aesthetic is modern and minimalist.

So, how can we implement  
**relationships** in a **relational**  
database?

Relational databases (in simple terms)

We use **keys** to **link** tables together

---

A table's **primary key** is used to link other tables back to **itself**. It is then called the **foreign key** in the other tables.

# LAB 1

# Within your team

- Work in pairs on your project
- Identify the entities (look for nouns)
- Identify the relationships between the entities
- Identify the important attributes for each entity
- As a team outline on paper, if not all, at least the 5 most important entities and their relationships





Shall we **share?**



# Within your team

- Is there something you can improve in your design?
- Why?



# **RELATIONAL DATABASES PART II**



What happens if a key ends up  
pointing to a row which has been  
**deleted?**

Referential integrity (violation of)

Then, we say that we have **broken**  
the **referential integrity** of our  
database

---

Referential integrity (violation of)

Then, we say that we have **broken**  
the **referential integrity** of our  
database

---

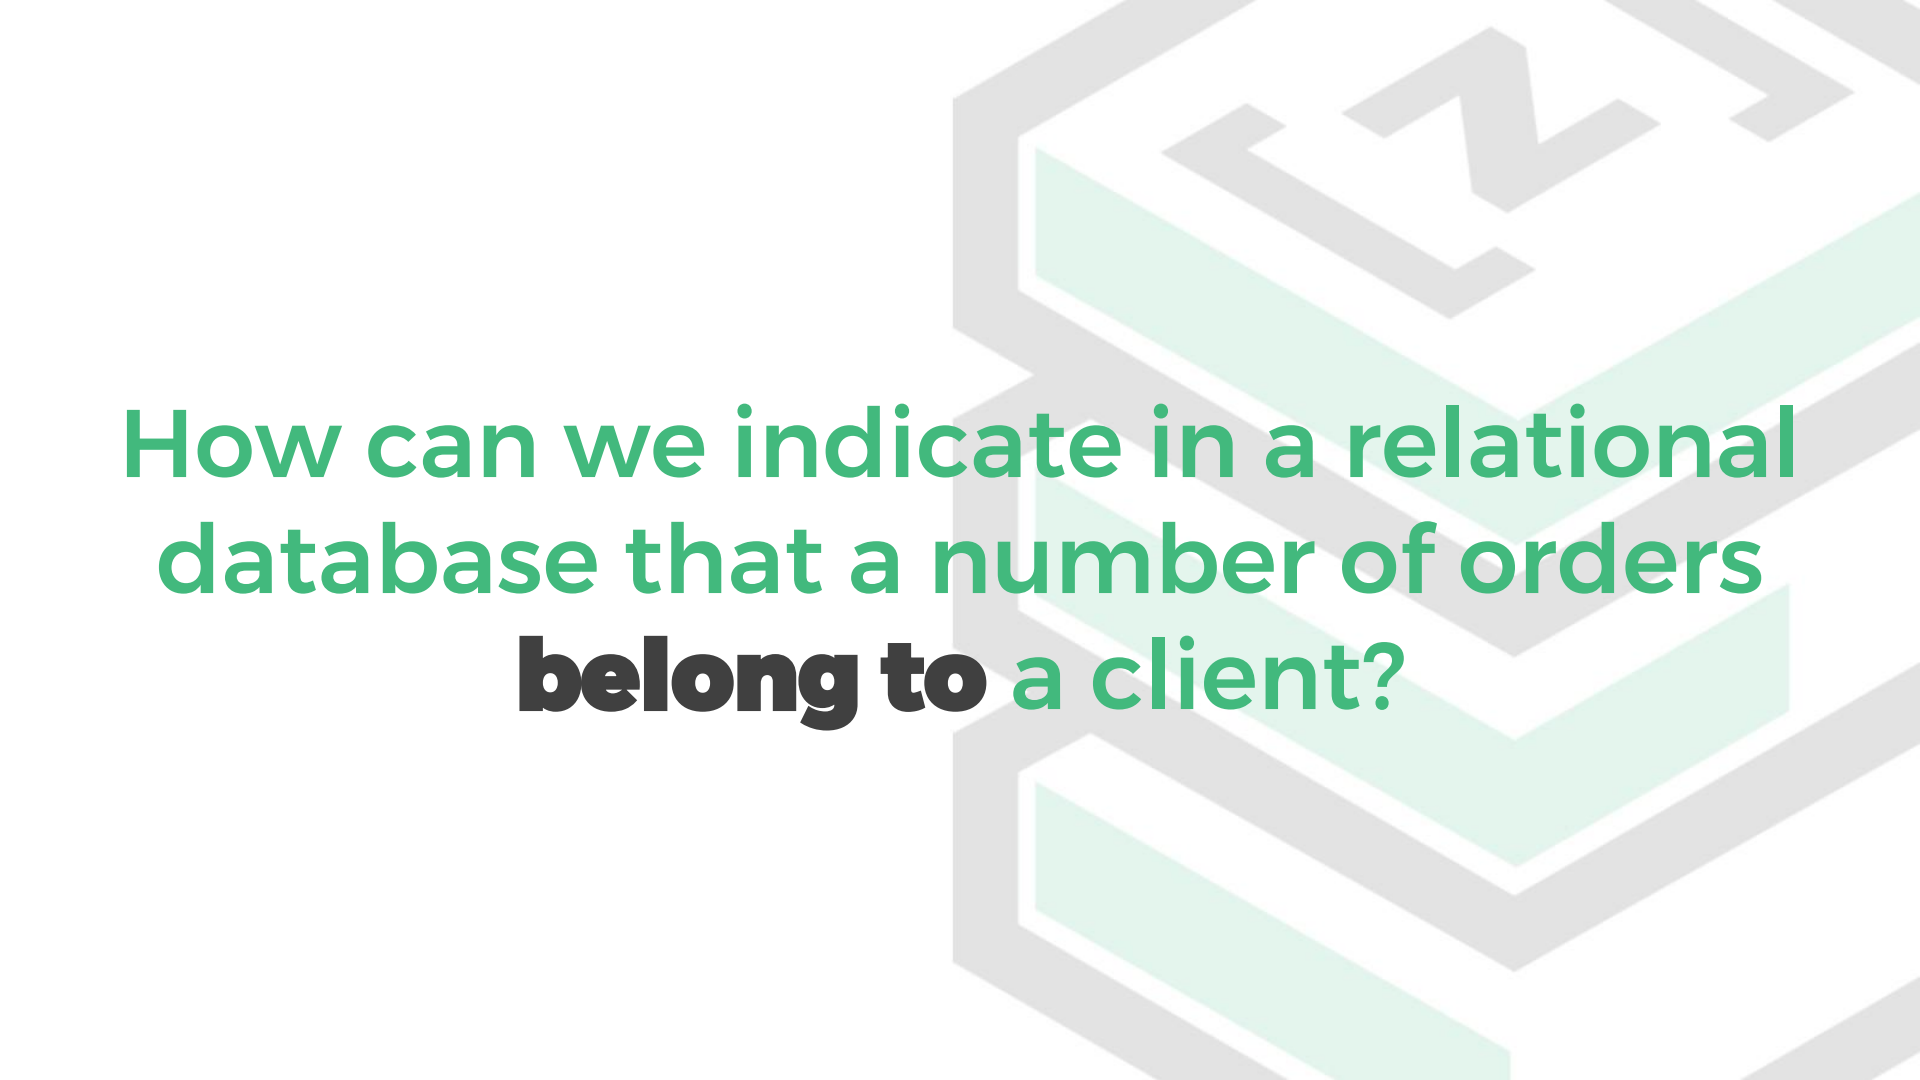
For referential integrity to hold, a **foreign key** must  
reference a valid, existing primary key in another table.

Referential integrity (violation of)

Then, we say that we have **broken**  
the **referential integrity** of our  
database

---

For referential integrity to hold, a **foreign key** must  
reference a valid existing primary key in another table.  
Relational databases enforce **constraints** to make sure  
these keys match up.



How can we indicate in a relational database that a number of orders **belong to** a client?



## Nature of relationships

- one to many
- many to one
- many to many
- one to one

Direct vs indirect relationships

Some of the relationships are just  
**derived** relationships and not  
**direct** relationships

---

In general, you don't have to do anything with derived relationships.

## Cardinality

Cardinality defines the **numbers** at each end of a relationship.

---

e.g., 1:N, 1:1, N:1, N:M



Shall we try an **example**?  
e.g., faculty - student - class



Why should we store the **foreign key** to another table rather than the data directly?



Shall we try another **example**?  
e.g., meetup group - members

Many to many

‘Many to many’ relationships **can not directly** be implemented in relational databases

---

We must break down the ‘many to many’ relationships into 2 ‘one to many’ and use an **intersecting entity**.



Shall we try yet another **example**?  
e.g., patient - medication





# **RELATIONAL DATABASES PART III**

Anomalies? Inconsistencies? Redundancy?

Normalisation is the **process** of improving the design of a database in order to achieve its **optimum** structure

Anomalies? Inconsistencies? Redundancy?

Normalisation tries to eliminate  
**redundancy**, improve **data**  
**integrity** and remove  
**inconsistencies**

Anomalies? Inconsistencies? Redundancy?

Normalisation tries to rid  
databases of **update, insertion**  
and **deletion** anomalies

## 1st Normal Form

Table has only data related to a single entity, no repeating attributes, has only **atomic values** and has a primary key

---

Querying and manipulating data within a table which is not normalised involves more **complexity** than is really needed.

# 1st Normal Form

**PRODUCT**

PRODUCT	COLOUR	PRICE
1	red, green	15.99
2	yellow	23.99
3	green	17.50
4	yellow, blue	9.99
5	red	29.99

1 NF

**PRODUCT PRICE**

PRODUCT	PRICE
1	15.99
2	23.99
3	17.50
4	9.99
5	29.99

**PRODUCT COLOUR**

PRODUCT	COLOUR
1	red
1	green
2	yellow
3	green
4	yellow
4	blue
5	red

The PRODUCT table is not in first normal form because the [COLOUR] column can contain multiple values. For example, the first row includes values "red" and "green."

## 2nd Normal Form

1NF + all non-PK attributes are fully **functionally dependent** on (the whole) of (every) PK

---

e.g., student\_email in student\_class table violates 2NF

## 2nd Normal Form

**PURCHASE DETAIL**

CUSTOMER ID	STORE ID	PURCHASE LOCATION
1	1	Los Angeles
1	3	San Francisco
2	1	Los Angeles
3	2	New York
4	3	San Francisco

2 NF

**PURCHASE**

CUSTOMER ID	STORE ID
1	1
1	3
2	1
3	2
4	3

**STORE**

STORE ID	PURCHASE LOCATION
1	Los Angeles
2	San Francisco
3	New York

The PURCHASE DETAIL table has a composite primary key [CUSTOMER ID, STORE ID]. The non-key attribute is [PURCHASE LOCATION]. In this case, [PURCHASE LOCATION] only depends on [STORE ID], which is only part of the primary key. Therefore, this table does not satisfy 2NF.



## 3rd Normal Form

2NF + Every non-prime attribute  
is **non-transitively dependent** on  
every key of the table

---

The only data that we store from another table should be just enough to create the relation. Otherwise it is redundant data. We can use the id to the other table and find all the info.

# 3rd Normal Form

**BOOK DETAIL**

BOOK ID	GENRE ID	GENRE TYPE	PRICE
1	1	Gardening	25.99
2	2	Sports	14.99
3	1	Gardening	10.00
4	3	Travel	12.99
5	2	Sports	17.99

3 NF

**BOOK**

BOOK ID	GENRE ID	PRICE
1	1	25.99
2	2	14.99
3	1	10.00
4	3	12.99
5	2	17.99

**GENRE**

GENRE ID	GENRE TYPE
1	Gardening
2	Sports
3	Travel

In the BOOK DETAIL table [BOOK ID] determines [GENRE ID], and [GENRE ID] determines [GENRE TYPE]. Therefore, [BOOK ID] determines [GENRE TYPE] via [GENRE ID] and we have transitive functional dependency, and this structure does not satisfy 3NF.



# **LAB**

## **2**

# Within your team

- Work in pairs on your project
- Identify the entities (look for nouns)
- Identify the relationships between the entities
- Transform 'many to many' relationships using intersecting entities
- Identify the attributes for each entity and normalise those entities to 3NF
- Identify primary keys
- As a team outline on paper your database design.




Shall we **share?**



# Within your team

- Is there something you can improve in your design?
- Why?



**SOME TIPS?**  
**(FROM BITTER EXPERIENCE)**

The background features a series of stacked, isometric blocks in light green and grey. A grey maze-like pattern is visible on the top surface of the uppermost block.

Where would you start your **data**  
**modelling** from?



The background features a series of stacked, isometric blocks in light green and white. To the right, there is a faint, grey maze-like pattern. The text is overlaid on the left side of the image.

Why would you choose a **relational** database?



**How powerful**  
(flexible, fast, accurate)  
do you want to be at storing,  
retrieving and producing  
information?

This material was based on

## **Mike Zellers - CISS143 Database Design course**

[https://www.youtube.com/watch?v=ZFpRzr8hCm4&index=12&list=PL38oYrcNWuOOUcLvQ\\_a7EEu8-mLqVfQZg](https://www.youtube.com/watch?v=ZFpRzr8hCm4&index=12&list=PL38oYrcNWuOOUcLvQ_a7EEu8-mLqVfQZg)

---

**Bitter experience**

**and** <http://www.1keydata.com/database-normalization/>

**“The goal of Ruby  
is to make  
programmers **happy**”**

- Yukihiro Matsumoto



A grayscale background image featuring several Stormtroopers from Star Wars. They are standing in a line, with some in the foreground and others further back. The setting appears to be an industrial or military facility with large structures and pipes in the background.

WE LOVE SHARING INSIGHTS  
**JOIN OUR NEWSLETTER**

**zanshin**labs.io

thank you